

Movie Dataset Analysis

This notebook performs an exploratory data analysis (EDA) on the TMDb 5000 movie and credits datasets. The goal is to understand the data, clean it, and extract useful insights about What factors influence movie revenue.

0. Installing Required Packages and Creating requirements.txt

```
In [ ]: import subprocess
import sys

# List of required packages
required_packages = ['pandas', 'numpy', 'matplotlib', 'seaborn', 'scikit-learn']

# Install packages
for package in required_packages:
    subprocess.check_call([sys.executable, "-m", "pip", "install", package])

# Create a requirements.txt file with current environment's packages
with open('requirements.txt', 'w') as f:
    subprocess.check_call([sys.executable, "-m", "pip", "freeze"], stdout=f)
```

1. Importing Necessary Libraries

We begin by importing the required libraries for data manipulation, visualization, and analysis.

```
In [ ]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Set visualization style
sns.set(style="darkgrid")
```

2. Loading the Data

Next, we load the datasets for movies and credits from CSV files.

```
In [ ]: # Load CSV files
credits_df = pd.read_csv('data/tmdb_5000_credits.csv')
movies_df = pd.read_csv('data/tmdb_5000_movies.csv')
```

Exploring Datasets

1. Display first few rows of each dataset to get an initial understanding of data.
2. Get a concise summary of datasets, including data types and non-null counts.
3. Check for missing values in each dataset.

```
In [ ]: # Display first few rows of each dataset
print("First Dataset:")
print(credits_df.head())

print("\nSecond Dataset:")
print(movies_df.head())

# Get summary of data
print(credits_df.info())
print(movies_df.info())

# Check for missing values
print(credits_df.isnull().sum())
print(movies_df.isnull().sum())
```

```
First Dataset:
  movie_id      title \
0      1995      Avatar
1      285  Pirates of the Caribbean: At World's End
2     206647      Spectre
3      49026  The Dark Knight Rises
4      49529    John Carter

  cast \
0 [{"cast_id": 242, "character": "Jake Sully", "...
1 [{"cast_id": 4, "character": "Captain Jack Spa...
2 [{"cast_id": 1, "character": "James Bond", "cr...
3 [{"cast_id": 2, "character": "Bruce Wayne / Ba...
4 [{"cast_id": 5, "character": "John Carter", "c...

  crew
0 [{"credit_id": "52fe48009251416c750aca23", "de...
1 [{"credit_id": "52fe4232c3a36847f800b579", "de...
2 [{"credit_id": "54805967c3a36829b5002c41", "de...
3 [{"credit_id": "52fe4781c3a36847f81398c3", "de...
4 [{"credit_id": "52fe479ac3a36847f813eaa3", "de...

Second Dataset:
  budget      genres \
0 237000000 [{"id": 28, "name": "Action"}, {"id": 12, "nam...
1 300000000 [{"id": 12, "name": "Adventure"}, {"id": 14, "...
2 245000000 [{"id": 28, "name": "Action"}, {"id": 12, "nam...
3 250000000 [{"id": 28, "name": "Action"}, {"id": 80, "nam...
4 260000000 [{"id": 28, "name": "Action"}, {"id": 12, "nam...

  homepage      id \
0      http://www.avatarmovie.com/      1995
1 http://disney.go.com/disneypictures/pirates/      285
2 http://www.sonypictures.com/movies/spectre/      206647
3      http://www.thedarkknightises.com/      49026
4      http://movies.disney.com/john-carter      49529

  keywords original_language \
0 [{"id": 1463, "name": "culture clash"}, {"id":...      en
1 [{"id": 270, "name": "ocean"}, {"id": 726, "na...      en
2 [{"id": 470, "name": "spy"}, {"id": 818, "name...      en
3 [{"id": 849, "name": "dc comics"}, {"id": 853,...      en
4 [{"id": 818, "name": "based on novel"}, {"id":...      en

  original_title \
0      Avatar
1  Pirates of the Caribbean: At World's End
2      Spectre
3      The Dark Knight Rises
4      John Carter

  overview popularity \
0 In the 22nd century, a paraplegic Marine is di... 150.437577
1 Captain Barbossa, long believed to be dead, ha... 139.082615
2 A cryptic message from Bond's past sends him o... 107.376788
3 Following the death of District Attorney Harve... 112.312950
4 John Carter is a war-weary, former military ca... 43.926995

  production_companies \
0 [{"name": "Ingenious Film Partners", "id": 289...
1 [{"name": "Walt Disney Pictures", "id": 2}, {"...
2 [{"name": "Columbia Pictures", "id": 5}, {"nam...
3 [{"name": "Legendary Pictures", "id": 923}, {"...
4 [{"name": "Walt Disney Pictures", "id": 2}]

  production_countries release_date      revenue \
0 [{"iso_3166_1": "US", "name": "United States o... 2009-12-10 2787965087
1 [{"iso_3166_1": "US", "name": "United States o... 2007-05-19 961000000
2 [{"iso_3166_1": "GB", "name": "United Kingdom"... 2015-10-26 880674609
3 [{"iso_3166_1": "US", "name": "United States o... 2012-07-16 1084939099
4 [{"iso_3166_1": "US", "name": "United States o... 2012-03-07 284139100

  runtime      spoken_languages      status \
0      162.0 [{"iso_639_1": "en", "name": "English"}, {"iso... Released
1      169.0 [{"iso_639_1": "en", "name": "English"}] Released
2      148.0 [{"iso_639_1": "fr", "name": "Fran\u00e7ais"},... Released
3      165.0 [{"iso_639_1": "en", "name": "English"}] Released
4      132.0 [{"iso_639_1": "en", "name": "English"}] Released

  tagline \
0      Enter the World of Pandora.
1 At the end of the world, the adventure begins.
2      A Plan No One Escapes
3      The Legend Ends
4      Lost in our world, found in another.

  title vote_average vote_count
0      Avatar          7.2      11800
1  Pirates of the Caribbean: At World's End          6.9      4500
2      Spectre          6.3      4466
3      The Dark Knight Rises          7.6      9106
4      John Carter          6.1      2124
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4803 entries, 0 to 4802
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   movie_id    4803 non-null   int64
 1   title       4803 non-null   object
 2   cast        4803 non-null   object
 3   crew        4803 non-null   object
dtypes: int64(1), object(3)
memory usage: 150.2+ KB
None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4803 entries, 0 to 4802
Data columns (total 20 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   budget              4803 non-null   int64
 1   genres              4803 non-null   object
 2   homepage            1712 non-null   object
 3   id                  4803 non-null   int64
 4   keywords            4803 non-null   object
 5   original_language   4803 non-null   object
 6   original_title      4803 non-null   object
 7   overview            4800 non-null   object
 8   popularity          4803 non-null   float64
 9   production_companies 4803 non-null   object
10   production_countries 4803 non-null   object
11   release_date        4802 non-null   object
12   revenue             4803 non-null   int64
13   runtime             4801 non-null   float64
14   spoken_languages    4803 non-null   object
15   status              4803 non-null   object
16   tagline             3959 non-null   object
17   title               4803 non-null   object
18   vote_average        4803 non-null   float64
19   vote_count          4803 non-null   int64
dtypes: float64(3), int64(4), object(13)
memory usage: 750.6+ KB
None
movie_id    0
title       0
cast        0
crew        0
dtype: int64
budget              0
genres              0
homepage            3091
id                  0
keywords            0
original_language   0
original_title      0
overview            3
popularity          0
production_companies 0
production_countries 0
release_date        1
revenue             0
runtime             2
spoken_languages    0
status              0
tagline            844
title               0
vote_average        0
vote_count          0
dtype: int64
```

4. Merging Datasets

To analyze the data together, we'll merge the movies and credits datasets based on the `movie_id` and `id` columns.

```
In [ ]: # Merge datasets on 'movie_id' and 'id'
combined_df = pd.merge(movies_df, credits_df, left_on='id', right_on='movie_id', how='inner')

# Display first few rows of combined dataset
print("Combined Dataset - First Few Rows:")
print(combined_df.head())

# Get summary of combined data
print("\nCombined Dataset - Summary:")
print(combined_df.info())
```

```
Combined Dataset - First Few Rows:
      budget      genres \
0  237000000 [{"id": 28, "name": "Action"}, {"id": 12, "nam...
1  300000000 [{"id": 12, "name": "Adventure"}, {"id": 14, "...
2  245000000 [{"id": 28, "name": "Action"}, {"id": 12, "nam...
3  250000000 [{"id": 28, "name": "Action"}, {"id": 80, "nam...
4  260000000 [{"id": 28, "name": "Action"}, {"id": 12, "nam...

      homepage      id \
0      http://www.avatarmovie.com/ 19995
1  http://disney.go.com/disneypictures/pirates/ 285
2      http://www.sonypictures.com/movies/spectre/ 206647
3      http://www.thedarkknighttrises.com/ 49026
4      http://movies.disney.com/john-carter 49529

      keywords original_language \
0 [{"id": 1463, "name": "culture clash"}, {"id":... en
1 [{"id": 270, "name": "ocean"}, {"id": 726, "na... en
2 [{"id": 470, "name": "spy"}, {"id": 818, "name... en
3 [{"id": 849, "name": "dc comics"}, {"id": 853,... en
4 [{"id": 818, "name": "based on novel"}, {"id":... en

      original_title \
0      Avatar
1  Pirates of the Caribbean: At World's End
2      Spectre
3      The Dark Knight Rises
4      John Carter

      overview popularity \
0  In the 22nd century, a paraplegic Marine is di... 150.437577
1  Captain Barbossa, long believed to be dead, ha... 139.082615
2  A cryptic message from Bond's past sends him o... 107.376788
3  Following the death of District Attorney Harve... 112.312950
4  John Carter is a war-weary, former military ca... 43.926995

      production_companies ... \
0 [{"name": "Ingenious Film Partners", "id": 289... ...
1 [{"name": "Walt Disney Pictures", "id": 2}, {""... ...
2 [{"name": "Columbia Pictures", "id": 5}, {"nam... ...
3 [{"name": "Legendary Pictures", "id": 923}, {""... ...
4      [{"name": "Walt Disney Pictures", "id": 2}] ...

      spoken_languages      status \
0 [{"iso_639_1": "en", "name": "English"}, {"iso... Released
1      [{"iso_639_1": "en", "name": "English"}] Released
2 [{"iso_639_1": "fr", "name": "Fran\u00e7ais"},... Released
3      [{"iso_639_1": "en", "name": "English"}] Released
4      [{"iso_639_1": "en", "name": "English"}] Released

      tagline \
0      Enter the World of Pandora.
1  At the end of the world, the adventure begins.
2      A Plan No One Escapes
3      The Legend Ends
4      Lost in our world, found in another.

      title_x vote_average vote_count movie_id \
0      Avatar 7.2 11800 19995
1  Pirates of the Caribbean: At World's End 6.9 4500 285
2      Spectre 6.3 4466 206647
3      The Dark Knight Rises 7.6 9106 49026
4      John Carter 6.1 2124 49529

      title_y \
0      Avatar
1  Pirates of the Caribbean: At World's End
2      Spectre
3      The Dark Knight Rises
4      John Carter

      cast \
0 [{"cast_id": 242, "character": "Jake Sully", "...
1 [{"cast_id": 4, "character": "Captain Jack Spa...
2 [{"cast_id": 1, "character": "James Bond", "cr...
3 [{"cast_id": 2, "character": "Bruce Wayne / Ba...
4 [{"cast_id": 5, "character": "John Carter", "c...

      crew
0 [{"credit_id": "52fe48009251416c750aca23", "de...
1 [{"credit_id": "52fe4232c3a36847f800b579", "de...
2 [{"credit_id": "54805967c3a36829b5002c41", "de...
3 [{"credit_id": "52fe4781c3a36847f81398c3", "de...
4 [{"credit_id": "52fe479ac3a36847f813eaa3", "de...

[5 rows x 24 columns]
```

```
Combined Dataset - Summary:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4803 entries, 0 to 4802
Data columns (total 24 columns):
#   Column      Non-Null Count  Dtype
---  -
#   Column      Non-Null Count  Dtype
```

```
0    budget      4803 non-null    int64
1    genres      4803 non-null    object
2    homepage    1712 non-null    object
3    id          4803 non-null    int64
4    keywords    4803 non-null    object
5    original_language  4803 non-null    object
6    original_title  4803 non-null    object
7    overview    4800 non-null    object
8    popularity  4803 non-null    float64
9    production_companies  4803 non-null    object
10   production_countries  4803 non-null    object
11   release_date  4802 non-null    object
12   revenue      4803 non-null    int64
13   runtime      4801 non-null    float64
14   spoken_languages  4803 non-null    object
15   status       4803 non-null    object
16   tagline      3959 non-null    object
17   title_x      4803 non-null    object
18   vote_average  4803 non-null    float64
19   vote_count   4803 non-null    int64
20   movie_id     4803 non-null    int64
21   title_y      4803 non-null    object
22   cast         4803 non-null    object
23   crew         4803 non-null    object
dtypes: float64(3), int64(5), object(16)
memory usage: 900.7+ KB
None
```


5. Visualizing Movie Budgets

We will now create a histogram to visualize the distribution of movie budgets.

```
In [ ]: plt.figure(figsize=(10, 6))
sns.histplot(combined_df['budget'], bins=30, kde=True)

# Customize x-axis Labels to show budget in millions
plt.gca().xaxis.set_major_formatter(plt.FuncFormatter(lambda x, _: f'${x/1e6:.0f}M'))

plt.title('Distribution of Movie Budgets')
plt.xlabel('Budget (in USD millions)')
plt.ylabel('Frequency')
plt.show()
```



6. Budget vs Revenue Scatter Plot

We will plot a scatter plot to visualize the relationship between movie budgets and revenues.

```
In [ ]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

plt.figure(figsize=(10, 10))

# Scatter plot with regression line
```

```
sns.regplot(
    data=combined_df,
    x='budget',
    y='revenue',
    scatter_kws={'s': 20, 'edgecolor': 'white', 'linewidths': 0.25}, # Use linewidths instead of linewidth
    line_kws={'color': 'red', 'linewidth': 1},
    ci=None
)

# Customize x-axis labels to show budget in millions
plt.gca().xaxis.set_major_formatter(plt.FuncFormatter(lambda x, _ : f'${x/1e6:.0f}M'))

# Customize y-axis labels to show revenue in billions
plt.gca().yaxis.set_major_formatter(plt.FuncFormatter(lambda y, _ : f'${y/1e9:.1f}B'))

# Calculate R2 value
X = combined_df['budget'].values.reshape(-1, 1)
y = combined_df['revenue'].values
model = LinearRegression().fit(X, y)
r2 = model.score(X, y)

# Annotate R2 value on plot
plt.text(0.05, 0.95, f'$R^2$: {r2:.2f}', transform=plt.gca().transAxes, fontsize=12, verticalalignment='top')

# Plot title and labels
plt.title('Budget vs Revenue with Regression Line')
plt.xlabel('Budget (in USD millions)')
plt.ylabel('Revenue (in USD billions)')
plt.show()
```



7. Analyzing Genres

Let's explore the genres in the dataset, compute average ratings by genre as well as correlation between genre and budget-to-revenue ratios, and visualize the results.

```
In [ ]: import json
from sklearn.linear_model import LinearRegression
```

```

# Parse 'genres' column from JSON strings to Python objects (List of dictionaries)
combined_df['genres_parsed'] = combined_df['genres'].apply(json.loads)

# Extract genre names directly from Lists of dictionaries
combined_df['genre_names'] = combined_df['genres_parsed'].apply(lambda x: [i['name'] for i in x])

# Explode genre names into separate rows
genre_df = combined_df.explode('genre_names')

# Calculate average rating per genre
avg_rating_by_genre = genre_df.groupby('genre_names')['vote_average'].mean().sort_values(ascending=False)

# Plot average rating by genre
plt.figure(figsize=(12, 8))
sns.barplot(x=avg_rating_by_genre, y=avg_rating_by_genre.index)
plt.title('Average Rating by Genre')
plt.xlabel('Average Rating')
plt.xlim(5.75, 7) # Set x-axis limits 5.75-7
plt.ylabel('Genre')
plt.show()

# --- Explore Correlation between Genre and Budget-to-Revenue Ratio ---

# Calculate budget-to-revenue ratio
genre_df['budget_to_revenue_ratio'] = genre_df['revenue'] / genre_df['budget']

# Remove rows with NaN, inf, or very large values in 'budget_to_revenue_ratio'
genre_df = genre_df.replace([np.inf, -np.inf], np.nan) # Replace inf values with NaN
genre_df = genre_df.dropna(subset=['budget_to_revenue_ratio', 'vote_average']) # Drop rows with NaN values
genre_df = genre_df[genre_df['budget_to_revenue_ratio'] < 1e10] # Filter out extremely large values

# Group by genre and calculate average budget-to-revenue ratio
avg_budget_to_revenue_by_genre = genre_df.groupby('genre_names')['budget_to_revenue_ratio'].mean().sort_values(ascending=False)

# Plot average budget-to-revenue ratio by genre with Logarithmic x-axis
plt.figure(figsize=(12, 8))
sns.barplot(x=avg_budget_to_revenue_by_genre, y=avg_budget_to_revenue_by_genre.index)
plt.xscale('log') # Apply log scale to x-axis
plt.title('Average Budget-to-Revenue Ratio by Genre (Logarithmic Scale)')
plt.xlabel('Budget-to-Revenue Ratio (Log Scale)')
plt.ylabel('Genre')
plt.show()

# --- Scatter Plot: Budget-to-Revenue Ratio vs. Movie Rating with Regression Line ---

# Calculate budget-to-revenue ratio for combined dataset
combined_df['budget_to_revenue_ratio'] = combined_df['revenue'] / combined_df['budget']

# Remove rows with NaN, 0, inf, or very large values in 'budget_to_revenue_ratio' or 'vote_average'
combined_df = combined_df.replace([np.inf, -np.inf], np.nan)
combined_df = combined_df.dropna(subset=['budget_to_revenue_ratio', 'vote_average'])
combined_df = combined_df[combined_df['budget_to_revenue_ratio'] < 1e10]
combined_df = combined_df[combined_df['vote_average'] > 0]

# --- Ensure X and y are from same filtered dataset ---

# Add a small constant to avoid log(0) issues
epsilon = 1e-10
combined_df['log_budget_to_revenue_ratio'] = np.log10(combined_df['budget_to_revenue_ratio'] + epsilon)

# Remove any infinite or NaN values that resulted from Log transformation
combined_df = combined_df.dropna(subset=['log_budget_to_revenue_ratio', 'vote_average'])

# Prepare data for regression analysis
X = combined_df[['vote_average']].values # Independent variable: movie ratings
y_log = combined_df['log_budget_to_revenue_ratio'].values # Dependent variable: Log-transformed budget-to-revenue ratio

# Initialize & Fit Model
model = LinearRegression()
model.fit(X, y_log)

# Predict and calculate R-squared for Log-transformed data
r2_log = model.score(X, y_log)
print(f'R² value for correlation between movie rating and log-transformed budget-to-revenue ratio: {r2_log:.2f}')

# Predict y values for Log-transformed B2R
y_log_pred = model.predict(X)

# Define threshold for filtering out small budget-to-revenue ratios
threshold = 1e-10

# Filter out data points where budget_to_revenue_ratio is too small and create a copy
filtered_df = combined_df[combined_df['budget_to_revenue_ratio'] > threshold].copy()

# Proceed with Log transformation
filtered_df['log_budget_to_revenue_ratio'] = np.log10(filtered_df['budget_to_revenue_ratio'] + epsilon)

# Fit model and make predictions
X = filtered_df[['vote_average']].values # Ensure this is a numpy array
y_log = filtered_df['log_budget_to_revenue_ratio'].values

model = LinearRegression()
model.fit(X, y_log)

```



```

# Predict using model
y_log_pred = model.predict(X)

# Plot data again
plt.figure(figsize=(10, 6))
plt.scatter(filtered_df['vote_average'], filtered_df['log_budget_to_revenue_ratio'], color='blue', edgecolor='white', linewidths=0.5, s=20,
plt.plot(filtered_df['vote_average'], y_log_pred, color='red', linewidth=1.5, label=f'Linear regression (R² = {r2_log:.2f})')
plt.title('Movie Rating vs Log-Transformed Budget-to-Revenue Ratio')
plt.xlabel('Movie Rating (Vote Average)')
plt.ylabel('Log(Budget-to-Revenue Ratio)')
plt.legend()
plt.show()

# Drop rows with 0 for budget and revenue
combined_df = combined_df[(combined_df['budget'] > 0) & (combined_df['revenue'] > 0)]

# Log-transform budget and revenue
combined_df['log_budget'] = np.log10(combined_df['budget'])
combined_df['log_revenue'] = np.log10(combined_df['revenue'])

# --- Scatter Plot: Movie Rating vs. Log-Transformed Budget with Regression Line ---

# Prepare data for regression analysis
X_rating_log_budget = combined_df[['vote_average']].values # Independent variable: movie ratings
y_log_budget = combined_df['log_budget'].values # Dependent variable: Log-transformed budget


# Initialize & Fit Model
model_rating_log_budget = LinearRegression()
model_rating_log_budget.fit(X_rating_log_budget, y_log_budget)


# Predict and calculate R-squared for rating vs. Log-transformed budget
r2_rating_log_budget = model_rating_log_budget.score(X_rating_log_budget, y_log_budget)
print(f'R² value for correlation between movie rating and log-transformed budget: {r2_rating_log_budget:.2f}')

# Predict y values for rating vs. Log-transformed budget
y_rating_log_budget_pred = model_rating_log_budget.predict(X_rating_log_budget)


# Scatter plot with regression line for rating vs. Log-transformed budget
plt.figure(figsize=(10, 6))
plt.scatter(X_rating_log_budget, y_log_budget, color='green', edgecolor='white', linewidths=0.5, s=20, label='Data points')
plt.plot(X_rating_log_budget, y_rating_log_budget_pred, color='red', linewidth=1.5, label=f'Linear regression (R² = {r2_rating_log_budget:.2f})')
plt.title('Movie Rating vs Log-Transformed Budget')
plt.xlabel('Movie Rating (Vote Average)')
plt.ylabel('Log(Budget)')
plt.legend()
plt.show()

```






R^2 value for correlation between movie rating and log-transformed budget-to-revenue ratio: 0.05



R^2 value for correlation between movie rating and log-transformed budget: 0.01

Movie Rating vs Log-Transformed Budget



8. Cleaning and Transforming Cast Data

Here, we will parse the `cast` column from the credits dataset, extract relevant information, and clean the data.

```
In [ ]: import json

# Parse JSON strings
def safe_parse(val):
    if isinstance(val, str):
        try:
            return json.loads(val)
        except (ValueError, json.JSONDecodeError):
            return None # Handle invalid JSON format
    elif isinstance(val, list):
        return val # Already a List, no need to parse
    return None # Handle any other cases

# Apply safe parsing function to 'cast' column
credits_df['cast'] = credits_df['cast'].apply(safe_parse)

# Function to extract desired fields
def extract_cast_info(cast_list):
    return [{'character': cast['character'], 'gender': cast['gender'], 'name': cast['name']} for cast in cast_list]

# Apply function to 'cast' column
credits_df['cast_extracted'] = credits_df['cast'].apply(extract_cast_info)

# Explode 'cast_extracted' column so each row contains a single actor's info
exploded_cast_df = credits_df[['title', 'cast_extracted']].explode('cast_extracted', ignore_index=True)

# Convert list of dictionaries into separate columns
exploded_cast_df = pd.concat([exploded_cast_df.drop('cast_extracted', axis=1), exploded_cast_df['cast_extracted'].apply(pd.Series)], axis=1)

# Drop column full of NaN values
exploded_cast_df = exploded_cast_df.drop(columns=[0])

# Change gender column to 'M' and 'F'
exploded_cast_df['gender'] = exploded_cast_df['gender'].replace({2: 'M', 1: 'F'})

# Display final DataFrame
print(exploded_cast_df.head())
```

	title	character	gender	name
0	Avatar	Jake Sully	M	Sam Worthington
1	Avatar	Neytiri	F	Zoe Saldana
2	Avatar	Dr. Grace Augustine	F	Sigourney Weaver
3	Avatar	Col. Quaritch	M	Stephen Lang
4	Avatar	Trudy Chacon	F	Michelle Rodriguez

9. Saving Cleaned Cast Data

We will now save the cleaned cast data to a new CSV file to reduce compute times for future analysis.

```
In [ ]: # Save new DataFrame to new CSV file
exploded_cast_df.to_csv('data/tmdb_5000_cast_cleaned.csv', index=False)
```

10. Merging Cleaned Cast Data with Movies Data

We will merge the cleaned cast data with the movies data to analyze actor performance metrics.

```
In [ ]: # Load clean cast data
cast_file_path = 'data/tmdb_5000_cast_cleaned.csv'
cast_df = pd.read_csv(cast_file_path)

# Merge datasets on 'title' column
combined_df = pd.merge(cast_df, movies_df[['id', 'title', 'budget', 'revenue']], on='title', how='inner')

# Display first few rows of combined DataFrame to confirm merge
print(combined_df.head())
```

	title	character	gender	name	id	budget	\
0	Avatar	Jake Sully	M	Sam Worthington	19995	237000000	
1	Avatar	Neytiri	F	Zoe Saldana	19995	237000000	
2	Avatar	Dr. Grace Augustine	F	Sigourney Weaver	19995	237000000	
3	Avatar	Col. Quaritch	M	Stephen Lang	19995	237000000	
4	Avatar	Trudy Chacon	F	Michelle Rodriguez	19995	237000000	
	revenue						
0	2787965087						
1	2787965087						
2	2787965087						
3	2787965087						
4	2787965087						

11. Actor Analysis: Budget-to-Revenue Ratio

We will calculate the budget-to-revenue ratio for movies and analyze the average ratios by actor.

```
In [ ]: # Calculate budget-to-revenue ratio
combined_df['budget_to_revenue_ratio'] = combined_df['revenue'] / combined_df['budget']

# Analyze average budget-to-revenue ratio by actor
actor_analysis = combined_df.groupby('name').agg({
    'budget': 'mean',
    'revenue': 'mean',
    'budget_to_revenue_ratio': 'mean'
}).reset_index()

# Sort results by budget-to-revenue ratio
actor_analysis = actor_analysis.sort_values(by='budget_to_revenue_ratio', ascending=False)

# Display top results
print(actor_analysis.head(10))
```

	name	budget	revenue	budget_to_revenue_ratio
54143	Zoë Hall	0.000000e+00	3.094813e+06	inf
35	Aakomon Jones	1.533333e+07	1.449227e+08	inf
37	Aamir Khan	1.100000e+06	6.015562e+06	inf
40	Aaron Abrams	3.325000e+07	6.390146e+07	inf
32498	Lydia Fox	0.000000e+00	8.646590e+05	inf
32483	Lutz Halbhubner	0.000000e+00	3.415310e+07	inf
54049	Zhang Ziyi	3.031778e+07	1.302567e+08	inf
54053	Zhao Hongfei	0.000000e+00	9.286394e+07	inf
54054	Zhao Wei	4.017050e+07	8.529568e+07	inf
23762	Jerry Hey	0.000000e+00	3.219520e+05	inf

12. Filtering Top Actors and Analyzing Movie Performances

We will focus on the top actors based on revenue-to-budget ratio, filter the movies involving them, and analyze their performance.

```
In [ ]: # Assuming original order of actors in 'cast_extracted' column indicates billing order,
# we'll assign a ranking based on order within each movie title.
cast_df['rank'] = cast_df.groupby('title').cumcount() + 1

# Filter to keep only first 5 actors (based on rank)
filtered_cast_df = cast_df[cast_df['rank'] <= 5]

# Merge filtered cast data with movies data on 'title' column
combined_df = pd.merge(filtered_cast_df, movies_df[['id', 'title', 'budget', 'revenue']], on='title', how='inner')

# Filter out rows where budget or revenue is zero
combined_df = combined_df[(combined_df['budget'] > 0) & (combined_df['revenue'] > 0)]

# Count number of movies per actor
```

```

actor_movie_counts = combined_df['name'].value_counts()

# Filter out actors who have appeared in fewer than 5 movies
actors_with_min_movies = actor_movie_counts[actor_movie_counts >= 5].index
filtered_combined_df = combined_df[combined_df['name'].isin(actors_with_min_movies)].copy()

# Calculate revenue-to-budget ratio for each movie
filtered_combined_df['revenue_to_budget_ratio'] = filtered_combined_df['revenue'] / filtered_combined_df['budget']

# Group by actor and calculate average revenue-to-budget ratio
actor_analysis = filtered_combined_df.groupby('name')['revenue_to_budget_ratio'].mean().reset_index()

# Sort results by revenue-to-budget ratio in descending order
actor_analysis = actor_analysis.sort_values(by='revenue_to_budget_ratio', ascending=False)

# Display top 10 actors by revenue-to-budget ratio
print(actor_analysis.head(10))

```

	name	revenue_to_budget_ratio
465	Kathleen Turner	125002.035637
671	Richard Dreyfuss	52.988923
222	Donald Pleasence	50.122596
350	Jamie Lee Curtis	33.708473
685	Robert Shaw	28.884410
359	Jason Mewes	25.396149
105	Carrie Fisher	20.689612
91	Bruce Campbell	17.489618
697	Roy Scheider	16.997388
645	Peter Coyote	16.792969

13. Analyzing Kathleen Turner Data

Kathleen Turner's `revenue_to_budget_ratio` was much higher than expected. Let's pull her data to find the cause.

```

In [ ]: # Filter combined DataFrame for Kathleen Turner's entries
kathleen_turner_df = filtered_combined_df[filtered_combined_df['name'] == 'Kathleen Turner']

# Display movies and their corresponding budgets and revenues
print(kathleen_turner_df[['title', 'budget', 'revenue', 'revenue_to_budget_ratio']])

```

	title	budget	revenue	revenue_to_budget_ratio
3448	Marley & Me	6000000	244082376	4.068040
5863	Dumb and Dumber To	40000000	169837010	4.245925
8677	A Simple Wish	28000000	8345056	0.298038
11668	Peggy Sue Got Married	18000000	41382841	2.299047
13724	Baby Geniuses	12000000	36450736	3.037561
13804	Serial Mom	13000000	7820688	0.601591
15668	Nurse 3-D	10	10000000	1000000.000000
17363	The Virgin Suicides	6000000	10409377	1.734896

14. Analyzing Top Actors by Revenue-to-Budget Ratio

We will continue by assigning a ranking to actors based on the order in which they appear in the `cast_extracted` column. Then, we will filter the top 5 actors for each movie and analyze the average revenue-to-budget ratio for these actors. We also need to eliminate placeholder in `budget` to avoid erroneous ratios.

```

In [ ]: # Assuming original order of actors in 'cast_extracted' column indicates billing order,
# we'll assign a ranking based on order within each movie title.
cast_df['rank'] = cast_df.groupby('title').cumcount() + 1

# Filter to keep only first 5 actors (based on rank)
filtered_cast_df = cast_df[cast_df['rank'] <= 5]

# Merge filtered cast data with movies data on 'title' column
combined_df = pd.merge(filtered_cast_df, movies_df[['id', 'title', 'budget', 'revenue']], on='title', how='inner')

# Filter out rows where budget or revenue is zero, and ensure budget exceeds $100,000
combined_df = combined_df[(combined_df['budget'] > 100000) & (combined_df['revenue'] > 0)]

# Calculate revenue-to-budget ratio for each movie
combined_df['revenue_to_budget_ratio'] = combined_df['revenue'] / combined_df['budget']

# Filter out actors who have appeared in fewer than 5 movies
actor_movie_counts = combined_df['name'].value_counts()
actors_with_min_movies = actor_movie_counts[actor_movie_counts >= 5].index
filtered_combined_df = combined_df[combined_df['name'].isin(actors_with_min_movies)]

# Verify if ratio calculation and filtering are correct
print(filtered_combined_df[['name', 'title', 'revenue_to_budget_ratio']].head())

# Group by actor and calculate average revenue-to-budget ratio
actor_analysis = filtered_combined_df.groupby('name')['revenue_to_budget_ratio'].mean().reset_index()

# Sort results by revenue-to-budget ratio in descending order
actor_analysis = actor_analysis.sort_values(by='revenue_to_budget_ratio', ascending=False)

```

```
# Display top 10 actors by revenue-to-budget ratio
print(actor_analysis.head(10))
```

	name	title \
0	Sam Worthington	Avatar
1	Zoe Saldana	Avatar
2	Sigourney Weaver	Avatar
4	Michelle Rodriguez	Avatar
5	Johnny Depp	Pirates of the Caribbean: At World's End

	revenue_to_budget_ratio
0	11.763566
1	11.763566
2	11.763566
4	11.763566
5	3.203333

	name	revenue_to_budget_ratio
671	Richard Dreyfuss	52.988923
222	Donald Pleasence	50.122596
350	Jamie Lee Curtis	33.708473
685	Robert Shaw	28.884410
105	Carrie Fisher	20.689612
91	Bruce Campbell	17.489618
676	Robert Carlyle	17.449983
696	Roy Scheider	16.997388
645	Peter Coyote	16.792969
15	Alec Guinness	16.679081

16. Further Analysis on Top Actors

We will further refine our analysis by focusing on the top 10 actors with the highest revenue-to-budget ratios. We will extract and display all movies involving these top actors to gain deeper insights.

```
In [ ]: # Assuming original order of actors in 'cast_extracted' column indicates billing order,
# we'll assign a ranking based on order within each movie title.
cast_df['rank'] = cast_df.groupby('title').cumcount() + 1

# Filter to keep only first 5 actors (based on rank)
filtered_cast_df = cast_df[cast_df['rank'] <= 5]

# Merge filtered cast data with movies data on 'title' column
combined_df = pd.merge(filtered_cast_df, movies_df[['id', 'title', 'budget', 'revenue']], on='title', how='inner')

# Filter out rows where budget or revenue is zero, and ensure budget exceeds $100,000
combined_df = combined_df[(combined_df['budget'] > 100000) & (combined_df['revenue'] > 0)]

# Calculate revenue-to-budget ratio for each movie
combined_df['revenue_to_budget_ratio'] = combined_df['revenue'] / combined_df['budget']

# Filter out actors who have appeared in fewer than 5 movies
actor_movie_counts = combined_df['name'].value_counts()
actors_with_min_movies = actor_movie_counts[actor_movie_counts >= 5].index
filtered_combined_df = combined_df[combined_df['name'].isin(actors_with_min_movies)]

# Group by actor and calculate average revenue-to-budget ratio
actor_analysis = filtered_combined_df.groupby('name')['revenue_to_budget_ratio'].mean().reset_index()

# Sort results by revenue-to-budget ratio in descending order to get top 10 actors
top_10_actors = actor_analysis.sort_values(by='revenue_to_budget_ratio', ascending=False).head(10)

# Extract names of top 10 actors
top_10_actor_names = top_10_actors['name'].tolist()

# Filter combined DataFrame for movies involving top 10 actors
top_10_movies = filtered_combined_df[filtered_combined_df['name'].isin(top_10_actor_names)]

# Sort by actor and then by movie title for clarity
top_10_movies_sorted = top_10_movies.sort_values(by=['name', 'title'])

# Display all movies and their ratios for top 10 actors
for actor in top_10_actor_names:
    print(f"\nActor: {actor}")
    actor_movies = top_10_movies_sorted[top_10_movies_sorted['name'] == actor]
    for index, row in actor_movies.iterrows():
        print(f"  Movie: {row['title']}, Revenue-to-Budget Ratio: {row['revenue_to_budget_ratio']:.2f}")
```

Actor: Richard Dreyfuss
Movie: American Graffiti, Revenue-to-Budget Ratio: 180.18
Movie: Close Encounters of the Third Kind, Revenue-to-Budget Ratio: 15.19
Movie: Jaws, Revenue-to-Budget Ratio: 67.24
Movie: My Life in Ruins, Revenue-to-Budget Ratio: 1.20
Movie: Poseidon, Revenue-to-Budget Ratio: 1.14

Actor: Donald Pleasence
Movie: Escape from New York, Revenue-to-Budget Ratio: 8.37
Movie: Halloween, Revenue-to-Budget Ratio: 233.33
Movie: Halloween 4: The Return of Michael Myers, Revenue-to-Budget Ratio: 3.55
Movie: Halloween 5: The Revenge of Michael Myers, Revenue-to-Budget Ratio: 2.33
Movie: Halloween: The Curse of Michael Myers, Revenue-to-Budget Ratio: 3.02

Actor: Jamie Lee Curtis
Movie: Drowning Mona, Revenue-to-Budget Ratio: 0.96
Movie: Freaky Friday, Revenue-to-Budget Ratio: 4.24
Movie: Halloween, Revenue-to-Budget Ratio: 233.33
Movie: Halloween: Resurrection, Revenue-to-Budget Ratio: 2.90
Movie: The Fog, Revenue-to-Budget Ratio: 21.38
Movie: The Tailor of Panama, Revenue-to-Budget Ratio: 1.33
Movie: Trading Places, Revenue-to-Budget Ratio: 2.23
Movie: True Lies, Revenue-to-Budget Ratio: 3.29

Actor: Robert Shaw
Movie: A Man for All Seasons, Revenue-to-Budget Ratio: 7.27
Movie: Force 10 from Navarone, Revenue-to-Budget Ratio: 1.45
Movie: From Russia with Love, Revenue-to-Budget Ratio: 39.45
Movie: Jaws, Revenue-to-Budget Ratio: 67.24
Movie: The Sting, Revenue-to-Budget Ratio: 29.02

Actor: Carrie Fisher
Movie: Return of the Jedi, Revenue-to-Budget Ratio: 17.70
Movie: Star Wars, Revenue-to-Budget Ratio: 70.49
Movie: The Empire Strikes Back, Revenue-to-Budget Ratio: 29.91
Movie: Under the Rainbow, Revenue-to-Budget Ratio: 0.11
Movie: Undiscovered, Revenue-to-Budget Ratio: 0.12
Movie: When Harry Met Sally..., Revenue-to-Budget Ratio: 5.80

Actor: Bruce Campbell
Movie: Evil Dead II, Revenue-to-Budget Ratio: 1.65
Movie: My Name Is Bruce, Revenue-to-Budget Ratio: 0.12
Movie: Serving Sara, Revenue-to-Budget Ratio: 0.58
Movie: The Ant Bully, Revenue-to-Budget Ratio: 1.10
Movie: The Evil Dead, Revenue-to-Budget Ratio: 84.00

Actor: Robert Carlyle
Movie: 28 Weeks Later, Revenue-to-Budget Ratio: 4.28
Movie: Eragon, Revenue-to-Budget Ratio: 2.49
Movie: The Full Monty, Revenue-to-Budget Ratio: 73.67
Movie: The World Is Not Enough, Revenue-to-Budget Ratio: 2.68
Movie: Trainspotting, Revenue-to-Budget Ratio: 4.12

Actor: Roy Scheider
Movie: Jaws, Revenue-to-Budget Ratio: 67.24
Movie: Jaws 2, Revenue-to-Budget Ratio: 9.39
Movie: Romeo Is Bleeding, Revenue-to-Budget Ratio: 0.28
Movie: Sorcerer, Revenue-to-Budget Ratio: 0.55
Movie: The French Connection, Revenue-to-Budget Ratio: 22.87
Movie: The Punisher, Revenue-to-Budget Ratio: 1.66

Actor: Peter Coyote
Movie: A Walk to Remember, Revenue-to-Budget Ratio: 3.75
Movie: E.T. the Extra-Terrestrial, Revenue-to-Budget Ratio: 75.52
Movie: Femme Fatale, Revenue-to-Budget Ratio: 0.48
Movie: Patch Adams, Revenue-to-Budget Ratio: 4.05
Movie: Sphere, Revenue-to-Budget Ratio: 0.17

Actor: Alec Guinness
Movie: A Passage to India, Revenue-to-Budget Ratio: 3.40
Movie: Doctor Zhivago, Revenue-to-Budget Ratio: 10.17
Movie: Lawrence of Arabia, Revenue-to-Budget Ratio: 4.67
Movie: Star Wars, Revenue-to-Budget Ratio: 70.49
Movie: The Bridge on the River Kwai, Revenue-to-Budget Ratio: 11.10
Movie: The Fall of the Roman Empire, Revenue-to-Budget Ratio: 0.25

17. Identifying and Analyzing the Top 10 Actors by Total Revenue-to-Budget Ratio

In this section, we aim to identify the top 10 actors based on their total revenue-to-budget ratio across all movies they have appeared in. The steps include:

- Assign Actor Ranking:** We assign a rank to actors within each movie based on their billing order (assuming the order in `cast_extracted` indicates billing order).
- Filter for Top 5 Actors per Movie:** We keep only the top 5 actors in each movie based on their rank.
- Merge with Movie Data:** We merge the filtered cast data with movie data to include information on budgets and revenues.
- Data Filtering:** We filter out movies where the budget or revenue is zero and ensure that the budget exceeds \$10,000.
- Highest Revenue Movie per Actor:** We identify the movie that generated the highest revenue for each actor.
- Calculate Total Revenue and Budget:** We calculate the total revenue and total budget for each actor across all movies.

7. **Revenue-to-Budget Ratio:** We compute the total revenue-to-budget ratio for each actor.
8. **Filtering Actors:** We filter out actors who have appeared in fewer than 5 movies to focus on those with substantial filmography.
9. **Merge with Character Information:** We merge the data with the highest revenue movie information to include character names in our analysis. This is not necessary but may aid viewers by associating the actor's character.
10. **Visualization:** We create a bar plot to visualize the top 10 actors by total revenue-to-budget ratio.

This visualization helps us understand which actors have been most profitable in terms of the revenue their movies have generated relative to the budgets.

```
In [ ]: # Assuming original order of actors in 'cast_extracted' column indicates billing order,
# we'll assign a ranking based on order within each movie title.
cast_df['rank'] = cast_df.groupby('title').cumcount() + 1

# Filter to keep only first 5 actors (based on rank)
filtered_cast_df = cast_df[cast_df['rank'] <= 5]

# Merge filtered cast data with movies data on 'title' column
combined_df = pd.merge(filtered_cast_df, movies_df[['id', 'title', 'budget', 'revenue']], on='title', how='inner')

# Filter out rows where budget or revenue is zero, and ensure budget exceeds $10,000
combined_df = combined_df[(combined_df['budget'] > 10000) & (combined_df['revenue'] > 0)]

# Identify highest revenue movie for each actor
idx = combined_df.groupby('name')['revenue'].idxmax()
highest_revenue_df = combined_df.loc[idx, ['name', 'title', 'revenue', 'budget', 'character']]

# Calculate total revenue and total budget for each actor
actor_totals = combined_df.groupby('name').agg(
    total_revenue=pd.NamedAgg(column='revenue', aggfunc='sum'),
    total_budget=pd.NamedAgg(column='budget', aggfunc='sum')
).reset_index()

# Calculate revenue-to-budget ratio for each actor
actor_totals['total_revenue_to_budget_ratio'] = actor_totals['total_revenue'] / actor_totals['total_budget']

# Filter out actors who have appeared in fewer than 5 movies
actor_movie_counts = combined_df['name'].value_counts()
actors_with_min_movies = actor_movie_counts[actor_movie_counts >= 5].index
actor_totals = actor_totals[actor_totals['name'].isin(actors_with_min_movies)]

# Merge with highest revenue movie info to include character names
actor_totals = pd.merge(actor_totals, highest_revenue_df[['name', 'character', 'title']], on='name')

# Combine actor name and character name for y-axis labels
actor_totals['name_with_character'] = actor_totals['name'] + ' \n(' + actor_totals['character'] + ')'

# Sort results by total revenue-to-budget ratio in descending order
actor_totals_sorted = actor_totals.sort_values(by='total_revenue_to_budget_ratio', ascending=False).head(10)

# Set up plot with a more manageable figure size
plt.figure(figsize=(14, 10), dpi=100)
ax = sns.barplot(x='total_revenue_to_budget_ratio', y='name_with_character', hue='name_with_character', data=actor_totals_sorted, palette='magma')

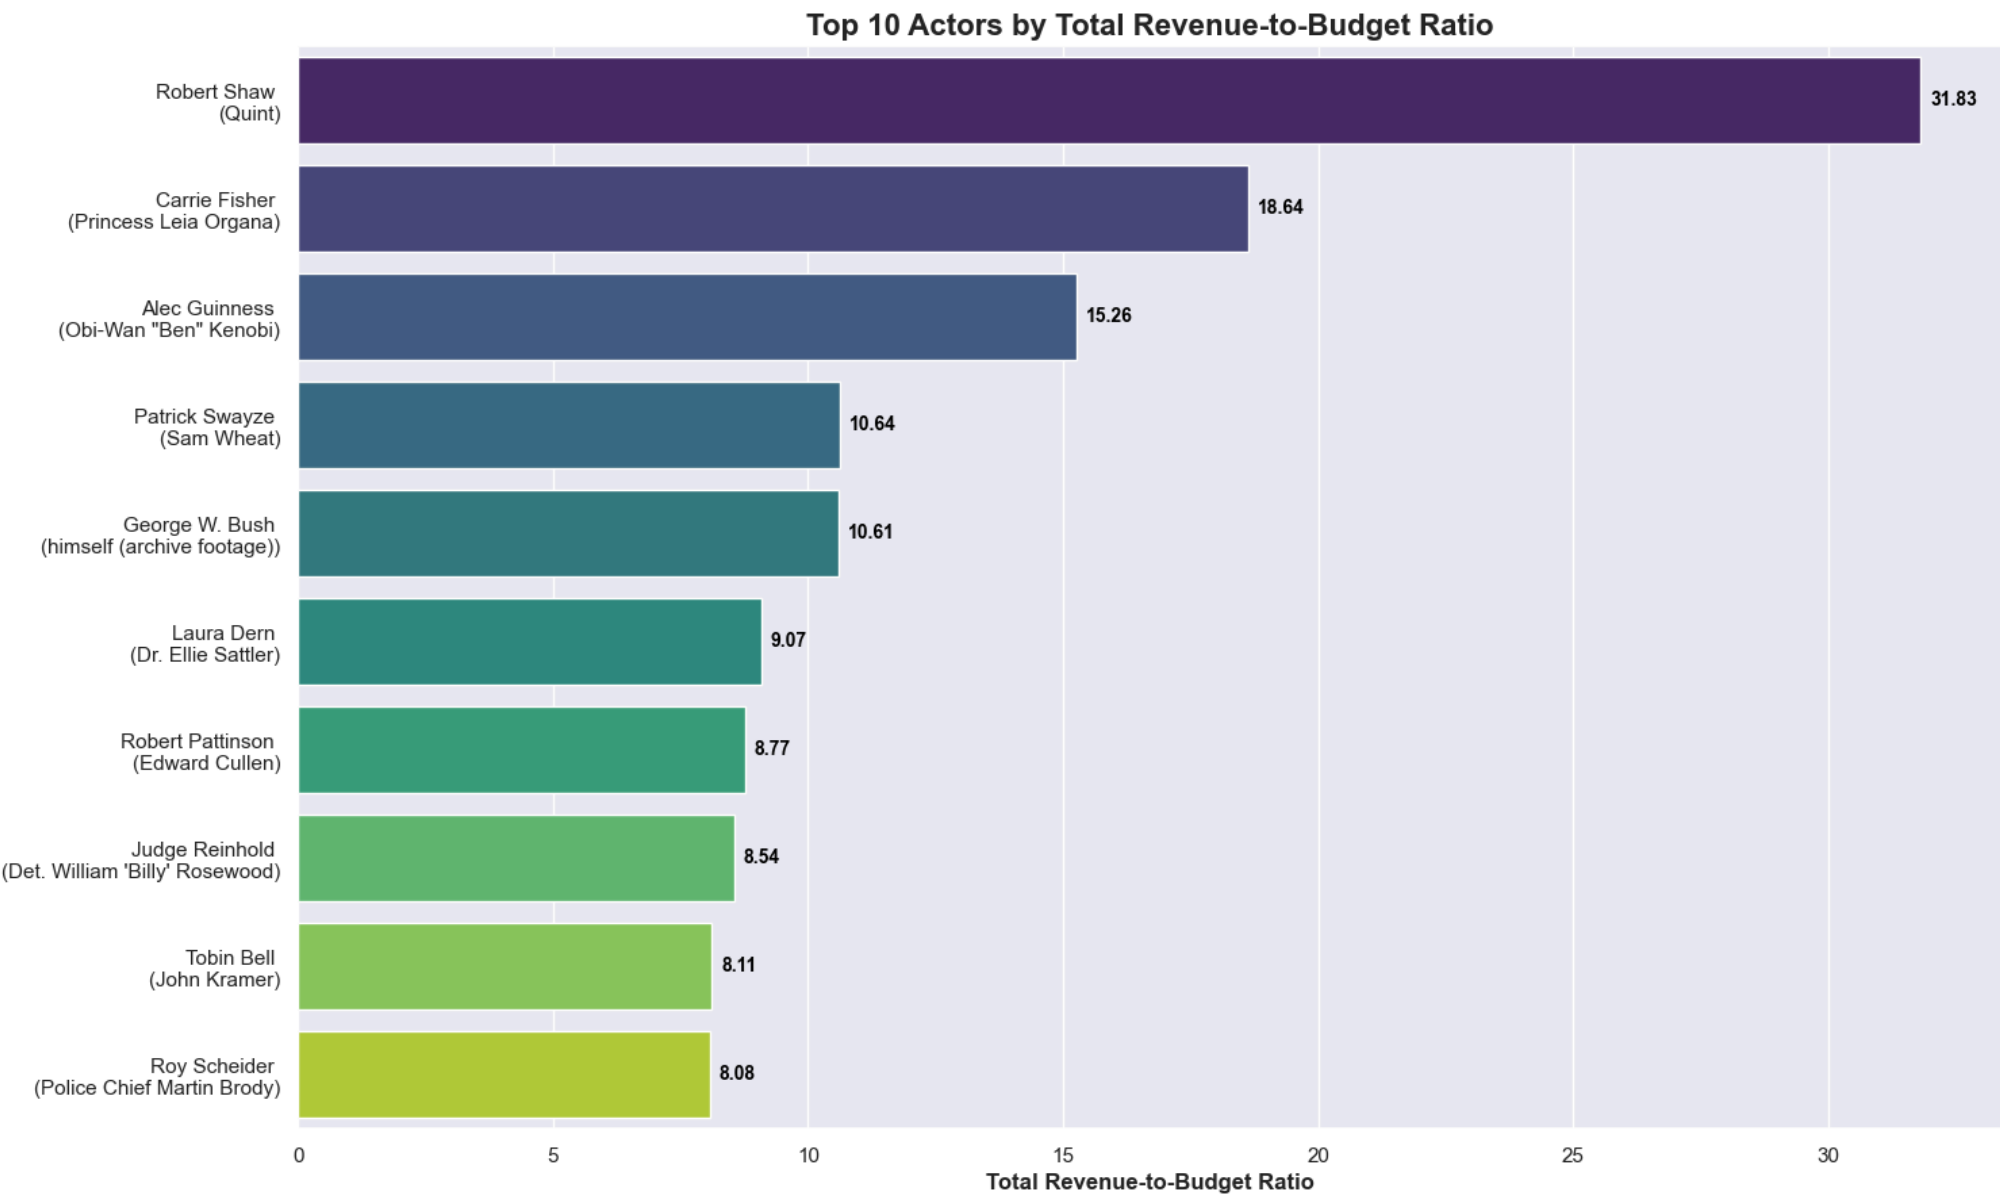
# Customize plot
plt.title('Top 10 Actors by Total Revenue-to-Budget Ratio', fontsize=16, weight='bold')
plt.xlabel('Total Revenue-to-Budget Ratio', fontsize=12, weight='bold')
ax.set_ylabel('')

# Use ax.bar_label for more reliable text annotation positioning
for container in ax.containers:
    ax.bar_label(container, fmt='%.2f', label_type='edge', fontsize=10, color='black', weight='bold', padding=5)

# Adjust layout manually to avoid clipping
plt.subplots_adjust(left=0.05, right=0.95, top=0.9, bottom=0.1)

# Show plot
plt.show()

# Display top 10 actors by total revenue-to-budget ratio and their highest revenue character
print(actor_totals_sorted[['name', 'total_revenue_to_budget_ratio', 'character', 'title']].head(10))
```

	name	total_revenue_to_budget_ratio	\
685	Robert Shaw	31.826884	
105	Carrie Fisher	18.642832	
15	Alec Guinness	15.261036	
631	Patrick Swayze	10.636271	
281	George W. Bush	10.606001	
502	Laura Dern	9.072935	
683	Robert Pattinson	8.767066	
445	Judge Reinhold	8.542249	
772	Tobin Bell	8.107859	
696	Roy Scheider	8.076311	

	character	title
685	Quint	Jaws
105	Princess Leia Organa	Star Wars
15	Obi-Wan "Ben" Kenobi	Star Wars
631	Sam Wheat	Ghost
281	himself (archive footage)	Fahrenheit 9/11
502	Dr. Ellie Sattler	Jurassic Park
683	Edward Cullen	The Twilight Saga: Breaking Dawn - Part 2
445	Det. William 'Billy' Rosewood	Beverly Hills Cop
772	John Kramer	Saw III
696	Police Chief Martin Brody	Jaws

18. Analyzing the Top 10 Movies by Revenue-to-Budget Ratio

In this section, we shift our focus from actors to movies. We aim to identify the top 10 movies based on their revenue-to-budget ratio. The steps include:

- Load and Filter Movie Data:** We load the movie dataset and filter out entries where the budget or revenue is zero, ensuring the budget exceeds \$10,000.
- Calculate Revenue-to-Budget Ratio:** For each movie, we compute the revenue-to-budget ratio.
- Identify Top 10 Movies:** We sort the movies by their revenue-to-budget ratio in descending order and select the top 10 movies.
- Visualization:** We create a bar plot to visualize these top 10 movies, allowing for a comparison of how efficiently each movie converted its budget into revenue.

This analysis is crucial to understanding which movies have been the most financially successful relative to their production costs.

```
In [ ]: # Load movies data
movies_df = pd.read_csv('data/tmdb_5000_movies.csv')

# Filter out rows where budget or revenue is zero, and ensure budget exceeds $10,000
filtered_movies_df = movies_df[(movies_df['budget'] > 10000) & (movies_df['revenue'] > 0)].copy()

# Calculate revenue-to-budget ratio for each movie using .loc to avoid SettingWithCopyWarning
filtered_movies_df.loc[:, 'revenue_to_budget_ratio'] = filtered_movies_df['revenue'] / filtered_movies_df['budget']

# Sort results by revenue-to-budget ratio in descending order and select top 10
top_movies_sorted = filtered_movies_df.sort_values(by='revenue_to_budget_ratio', ascending=False).head(10)


# Set up plot
plt.figure(figsize=(10, 6))
ax = sns.barplot(x='revenue_to_budget_ratio', y='title', hue='title', data=top_movies_sorted[['title', 'revenue_to_budget_ratio']], palette=
```

```
# Customize plot
plt.title('Top 10 Movies by Revenue-to-Budget Ratio (Budget > $10,000)', fontsize=16)
plt.xlabel('Revenue-to-Budget Ratio', fontsize=14)
ax.set_ylabel('')
plt.xscale('log') # Log scale to better visualize differences

# Annotate bars with exact ratio values
for index, value in enumerate(top_movies_sorted['revenue_to_budget_ratio']):
    plt.text(value, index, f'{value:.2f}', va='center')

# Show plot
plt.tight_layout()
plt.show()

# Display top 10 movies by revenue-to-budget ratio
print(top_movies_sorted[['title', 'revenue_to_budget_ratio']].head(10))
```



	title	revenue_to_budget_ratio
4577	Paranormal Activity	12890.386667
4496	The Blair Witch Project	4133.333333
4788	Pink Flamingos	500.000000
4742	Super Size Me	439.616585
4723	The Gallows	426.644100
4514	Open Water	420.522723
3159	The Texas Chain Saw Massacre	363.047059
4441	Bambi	311.709965
4668	The Stewardesses	270.000000
3737	Night of the Living Dead	263.157895

19. Identifying and Visualizing the Top 5 Most Profitable Movies

In this section, we focus on identifying the top 5 most profitable movies, where profitability is defined as the difference between the revenue and the budget. The steps involved are:

1. **Prepare Unique Movie Data:**

- We start by ensuring the `combined_df` is prepared with the required columns: `title`, `budget`, and `revenue`.
- We drop any duplicate movie titles, keeping only the entry with the highest revenue for each title.

2. **Calculate Profit:**

- We convert the `revenue` and `budget` columns to `float64` for accurate calculations.
- We compute the profit for each movie, converting the result into billions of dollars and rounding it to two decimal places.

3. **Identify Top 5 Profitable Movies:**

- We sort the movies by profit in descending order and select the top 5.

4. **Visualize with a Pie Chart:**

- We create a pie chart to visualize the distribution of profit among the top 5 movies.
- The pie chart includes labels for each movie title and displays the profit as a percentage of the total profit, along with the exact profit value in billions of dollars.

5. **Display the Results:**

- We display the names and profits of the top 5 most profitable movies.

This analysis highlights which movies were the most financially successful in terms of absolute profit, providing insights into the most lucrative titles in the dataset.

```
In [ ]: # Assuming combined_df is already prepared and contains 'title', 'budget', and 'revenue'
# Drop duplicate titles, keeping entry with highest revenue
unique_movies_df = combined_df.drop_duplicates(subset='title', keep='first').copy()

# Set column type before calculation
unique_movies_df = unique_movies_df.astype({'revenue': 'float64', 'budget': 'float64'})

# Calculate profit for each unique movie, convert to billions, float 2 decimals
unique_movies_df['profit'] = ((unique_movies_df['revenue'] - unique_movies_df['budget']) / 1e9).round(2)

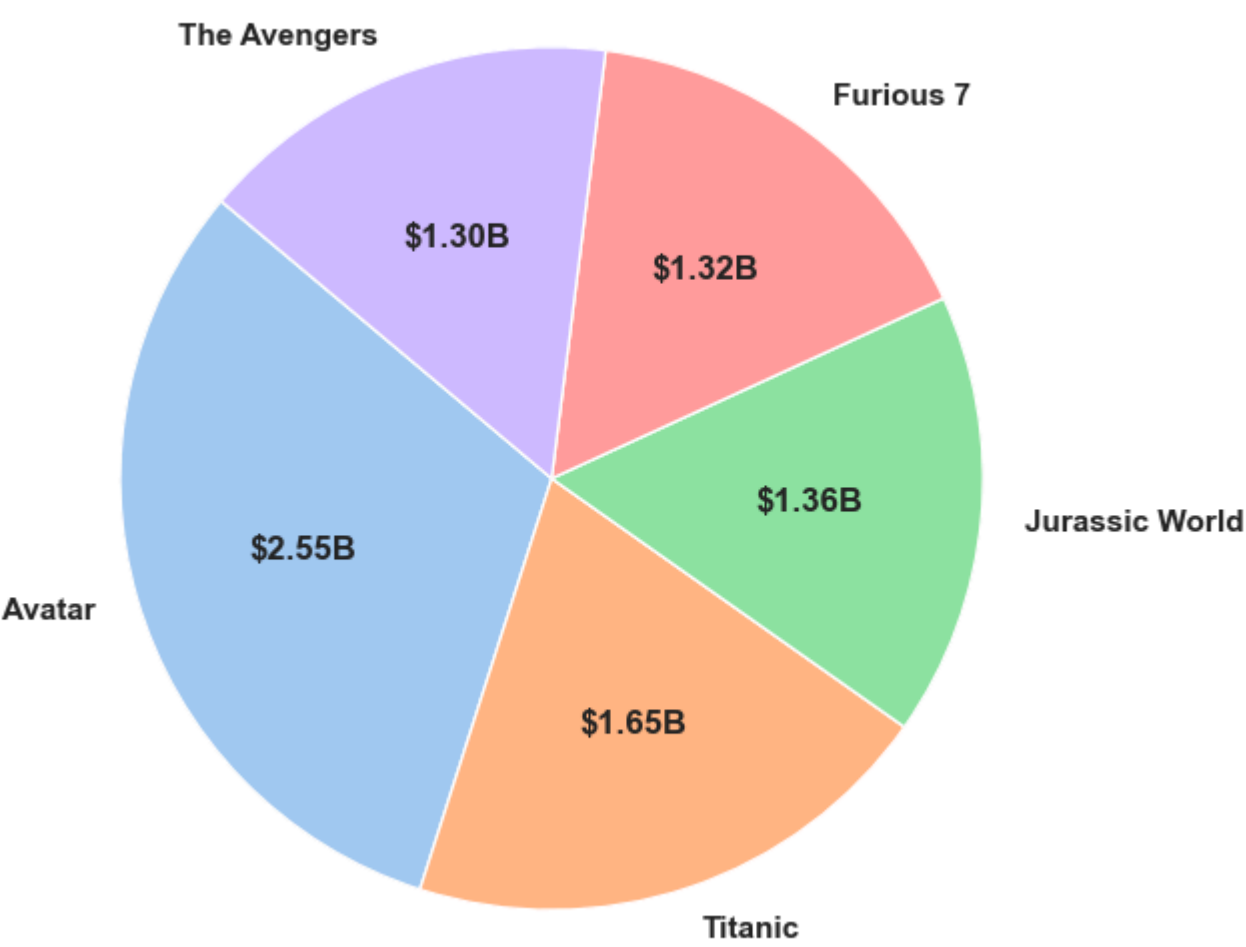
# Sort by profit and select top 5 most profitable movies
top_profitable_movies = unique_movies_df.sort_values(by='profit', ascending=False).head(5)

# Prepare data for pie chart
labels = top_profitable_movies['title']
sizes = top_profitable_movies['profit']

# Create pie chart
plt.figure(figsize=(10, 7))
plt.pie(sizes, labels=labels, autopct=lambda p: f'${p * sum(sizes) / 100:.2f}B', startangle=140,
        colors=sns.color_palette('pastel', len(labels)), textprops={'weight': 'bold'})
plt.title('Top 5 Most Profitable Movies (Revenue - Budget)', fontsize=16, weight='bold')
plt.show()

# Display top 5 most profitable movies with profit values in billions
formatted_output = top_profitable_movies.apply(lambda x: f'{x["title"]}: ${x["profit"]:.2f}B', axis=1)
print(formatted_output)
```

Top 5 Most Profitable Movies (Revenue - Budget)



```
0          Avatar: $2.55B
125         Titanic: $1.65B
140    Jurassic World: $1.36B
220          Furious 7: $1.32B
80          The Avengers: $1.30B
dtype: object
```